

# Data Transformation

Kylie Ariel Bemis

9/18/2018

# Introduction to Data Transformation

In order to visualize and (later) model data, we typically need to perform data transformation.

- ▶ Subset data by rows (observations) or columns (variables)
- ▶ Reorder the rows or columns of the data
- ▶ Calculate new variables from existing variables
- ▶ Calculate summary statistics of variables

One or more of these steps are often necessary to appropriately visualize or model a dataset.

# Introduction to Data Transformation (cont'd)

This week, we will discuss data transformations on **tidy** data. That is, data that is already in a form where:

- ▶ Each variable has its own column
- ▶ Each observation has its own row
- ▶ Each value has its own cell

*Next week*, we will discuss more on tidy data and how to perform the data wrangling necessary to clean up *untidy* data and turn it into tidy data.

# dplyr: A grammar of data manipulation

Provides a powerful, flexible “grammar” of data manipulation.

- ▶ Identify the most important data manipulation verbs and make them easy to use from R
- ▶ Provide fast performance for in-memory data with under-the-hood C++ implementation
- ▶ Use the same interface whether the data is stored in-memory or in a database on disk

Part of the “tidyverse” with ggplot2 and others.

# Why use dplyr?

The dplyr package (mostly) doesn't implement any functionality that is missing or impossible to perform in base R.

However, code written with dplyr:

- ▶ Can be more expressive, concise, and human-readable
- ▶ Can be more explicit in your intentions for data manipulation
- ▶ Can be faster (sometimes) due to C++ implementation
- ▶ Can also be used with databases on disk
- ▶ Integrates with other tidyverse functions and packages

I will also show you how to perform the same functionality using base R.

# Review of subsetting `data.frames` in R

- ▶ Simplifying (returns a vector)
  - ▶ Access individual columns using `df$name`
  - ▶ Access individual columns using `df[["name"]]`
- ▶ Preserving (returns a `data.frame`)
  - ▶ Subset rows using `df[i,]`
  - ▶ Subset rows using `df[c("name1", "name2"),]`
  - ▶ Subset columns using `df[,j]`
  - ▶ Subset columns using `df[,c("name1", "name2")]`
- ▶ You can subset both rows and columns at the same time
- ▶ Subsetting by a single column is always simplifying for `data.frames`
  - ▶ Change this behavior with `drop = FALSE`

```
df <- data.frame(x=c(1L, 2L, 5L, 9L),  
                 y=c('a', 'b', 'c', 'd'),  
                 `z !`=c(1.11, 2.22, 3.33, 4.0),  
                 row.names=c("Jo", "Ha", "Q", "Final"),  
                 check.names=FALSE,  
                 stringsAsFactors=FALSE)  
df
```

```
##      x y  z !  
## Jo   1 a 1.11  
## Ha   2 b 2.22  
## Q    5 c 3.33  
## Final 9 d 4.00
```

```
df$x
```

```
## [1] 1 2 5 9
```

```
df[["y"]]
```

```
## [1] "a" "b" "c" "d"
```



```
df$`z !`
```

```
## [1] 1.11 2.22 3.33 4.00
```

```
df[["z !"]]
```

```
## [1] 1.11 2.22 3.33 4.00
```

```
df[1:3,]
```

```
##      x y  z !  
## Jo  1 a 1.11  
## Ha  2 b 2.22  
## Q   5 c 3.33
```

```
df[c("Jo", "Ha", "Q"),]
```

```
##      x y  z !  
## Jo  1 a 1.11  
## Ha  2 b 2.22  
## Q   5 c 3.33
```

```
df[,2:3]
```

```
##           y  z !  
## Jo      a 1.11  
## Ha      b 2.22  
## Q       c 3.33  
## Final d 4.00
```

```
df[,c("y", "z !")]
```

```
##           y  z !  
## Jo      a 1.11  
## Ha      b 2.22  
## Q       c 3.33  
## Final d 4.00
```

```
df[, "z !"]
```

```
## [1] 1.11 2.22 3.33 4.00
```

```
df[, "z !", drop=FALSE]
```

```
##           z !  
## Jo      1.11  
## Ha      2.22  
## Q       3.33  
## Final   4.00
```

```
df[1:3, c("y", "z !")]
```

```
##      y  z !  
## Jo a 1.11  
## Ha b 2.22  
## Q  c 3.33
```

# Verbs in dplyr

Provides most commonly used data manipulation actions.

- ▶ `filter()` subsets data by rows/observations
- ▶ `arrange()` reorders data by rows/observations
- ▶ `select()` subsets data by columns/variables
- ▶ `mutate()` creates new columns/variables
- ▶ `summarise()` calculates summary statistics

Each can be applied over levels of a categorical variable with `group_by()`.

Each takes a `data.frame` as the first argument and outputs a new `data.frame`.

# Loading dplyr

```
library(dplyr)
```

```
##
```

```
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':
```

```
##
```

```
##      filter, lag
```

```
## The following objects are masked from 'package:base':
```

```
##
```

```
##      intersect, setdiff, setequal, union
```

# Name masking and namespaces in R

Several functions are “masked” when `dplyr` is loaded.

They can still be accessed by fully qualifying their names:

- ▶ `stats::filter()`
- ▶ `stats::lag()`



## Name masking and namespaces in R (cont'd)

Each package in R creates its own namespace.

R finds functions based on the order packages are loaded.

You can see this search path with `search()`:

```
search()
```

```
## [1] ".GlobalEnv"          "package:dplyr"        "package:stats"
## [4] "package:graphics"    "package:grDevices"    "package:utils"
## [7] "package:datasets"    "package:methods"      "Autoloads"
## [10] "package:base"
```

## Name masking and namespaces in R (cont'd)

When name conflicts occur, a warning about masked names is given.

You can always use `package::function()` to find the right one.

## Example dataset

Today we will explore the `flights` dataset also used in the homework and the **R4DS** book.

```
library(nycflights13)  
flights
```

Note that `flights` is actually a `tibble`, which is simply a special type of `data.frame`.

I will use `data.frame` to refer to both interchangeably. We will discuss the differences in more depth later in the course.

## Subsetting rows with filter()

Get only flights from October.

In base R:

```
flights[flights$month == 10,]
```

In dplyr:

```
filter(flights, month == 10)
```

```
filter(flights, month == 10)
```

```
## # A tibble: 28,889 x 19
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1  2013    10     1     447           500         -13     614
## 2  2013    10     1     522           517          5     735
## 3  2013    10     1     536           545         -9     809
## 4  2013    10     1     539           545         -6     801
## 5  2013    10     1     539           545         -6     917
## 6  2013    10     1     544           550         -6     912
## 7  2013    10     1     549           600        -11     653
## 8  2013    10     1     550           600        -10     648
## 9  2013    10     1     550           600        -10     649
## 10 2013    10     1     551           600         -9     727
```

```
## # ... with 28,879 more rows, and 12 more variables: sched_arr_time <
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <
## #   minute <dbl>, time_hour <dtm>
```

## Why flights\$month vs month?

Like ggplot2, dplyr uses non-standard evaluation to facilitate ease of interactive programming.

In base R:

```
flights[flights$month == 10,]
```

In dplyr:

```
filter(flights, month == 10)
```

This approach is useful for interactive data analysis, but can cause problems when used within user-defined functions.

We may discuss more on non-standard evaluation later in the course if there is interest.

# Anatomy of `filter()`

```
filter(data, condition1, condition2, condition3, ...)
```

- ▶ The first argument is the data
- ▶ The following arguments are vectorized logical expressions
- ▶ Additional arguments are joined by `&` (AND)
- ▶ Rows that evaluate to `TRUE` are kept
- ▶ Rows that evaluate to `FALSE` or `NA` are dropped

# Review of logical operators in R

- ▶ Standard comparison operators: `==`, `!=`, `>`, `<`, `>=`, `<=`.
  - ▶ Remember to use `==` instead of `=` when doing comparisons
- ▶ `dplyr::near()` for checking floating point equality
- ▶ `dplyr::between()` is a synonym for `a <= x & x <= b`
- ▶ `&`, `|`, and `!` are vectorized AND, OR, and NOT
  - ▶ Non-vectorized versions (for `if` statements) are `&&` and `||`
- ▶ `%in%` checks if an element exists in a set
  - ▶ E.g., `x %in% c(a,b)` is equivalent to `x == a | x == b`
- ▶ `is.na` to check for missing values
  - ▶ Remember that `NA == NA` evaluates to `NA`



Get only flights from Alaska Airlines or Hawaiian Airlines.

In base R:

```
flights[flights$carrier %in% c("AS", "HA"),]
```

In dplyr:

```
filter(flights, carrier %in% c("AS", "HA"))
```

```
filter(flights, carrier %in% c("AS", "HA"))
```

```
## # A tibble: 1,056 x 19
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1  2013     1     1     724           725          -1    1020
## 2  2013     1     1     857           900          -3    1516
## 3  2013     1     1    1808          1815          -7    2111
## 4  2013     1     2     722           725          -3     949
## 5  2013     1     2     909           900           9    1525
## 6  2013     1     2    1818          1815           3    2131
## 7  2013     1     3     724           725          -1    1012
## 8  2013     1     3     914           900          14    1504
## 9  2013     1     3    1817          1815           2    2121
## 10 2013     1     4     725           725           0    1031
```

```
## # ... with 1,046 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>,
## #   minute <dbl>, time_hour <dtm>
```

Get only flights from Alaska Airlines or Hawaiian Airlines.

In base R:

```
flights[flights$carrier == "AS" | flights$carrier == "HA",]
```

In dplyr:

```
filter(flights, carrier == "AS" | carrier == "HA")
```

```
filter(flights, carrier == "AS" | carrier == "HA")
```

```
## # A tibble: 1,056 x 19
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1  2013     1     1     724             725          -1    1020
## 2  2013     1     1     857             900          -3    1516
## 3  2013     1     1    1808            1815          -7    2111
## 4  2013     1     2     722             725          -3     949
## 5  2013     1     2     909             900           9    1525
## 6  2013     1     2    1818            1815           3    2131
## 7  2013     1     3     724             725          -1    1012
## 8  2013     1     3     914             900          14    1504
## 9  2013     1     3    1817            1815           2    2121
## 10 2013     1     4     725             725           0    1031
```

```
## # ... with 1,046 more rows, and 12 more variables: sched_arr_time <int>,
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <int>,
## #   minute <dbl>, time_hour <dtm>
```

Get only flights between Honolulu and JFK.

In base R:

```
flights[flights$origin == "JFK" & flights$dest == "HNL",]
```

In dplyr:

```
filter(flights, origin == "JFK" & dest == "HNL")
```

```
filter(flights, origin == "JFK" & dest == "HNL")
```

```
## # A tibble: 342 x 19
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1  2013     1     1     857           900          -3    1516
## 2  2013     1     2     909           900           9    1525
## 3  2013     1     3     914           900          14    1504
## 4  2013     1     4     900           900           0    1516
## 5  2013     1     5     858           900          -2    1519
## 6  2013     1     6    1019           900          79    1558
## 7  2013     1     7    1042           900         102    1620
## 8  2013     1     8     901           900           1    1504
## 9  2013     1     9     641           900        1301    1242
## 10 2013     1    10     859           900          -1    1449
```

```
## # ... with 332 more rows, and 12 more variables: sched_arr_time <int>
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <
## #   minute <dbl>, time_hour <dtm>
```

Get only flights between Honolulu and JFK.

In base R:

```
flights[flights$origin == "JFK" & flights$dest == "HNL",]
```

In dplyr:

```
filter(flights, origin == "JFK", dest == "HNL")
```

```
filter(flights, origin == "JFK", dest == "HNL")
```

```
## # A tibble: 342 x 19
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1  2013     1     1     857           900          -3    1516
## 2  2013     1     2     909           900           9    1525
## 3  2013     1     3     914           900          14    1504
## 4  2013     1     4     900           900           0    1516
## 5  2013     1     5     858           900          -2    1519
## 6  2013     1     6    1019           900          79    1558
## 7  2013     1     7    1042           900         102    1620
## 8  2013     1     8     901           900           1    1504
## 9  2013     1     9     641           900        1301    1242
## 10 2013     1    10     859           900          -1    1449
```

```
## # ... with 332 more rows, and 12 more variables: sched_arr_time <int>
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <
## #   minute <dbl>, time_hour <dtm>
```



## Reordering rows with `arrange()`

Sort by flights that departed with least delay (most ahead of schedule):

In base R:

```
flights[order(flights$dep_delay),]
```

In dplyr:

```
arrange(flights, dep_delay)
```

```
arrange(flights, dep_delay)
```

```
## # A tibble: 336,776 x 19
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1  2013    12     7     2040           2123         -43       40
## 2  2013     2     3     2022           2055         -33     2240
## 3  2013    11    10     1408           1440         -32    1549
## 4  2013     1    11     1900           1930         -30    2233
## 5  2013     1    29     1703           1730         -27    1947
## 6  2013     8     9      729           755         -26    1002
## 7  2013    10    23     1907           1932         -25    2143
## 8  2013     3    30     2030           2055         -25    2213
## 9  2013     3     2     1431           1455         -24    1601
## 10 2013     5     5      934           958         -24    1225
```

```
## # ... with 336,766 more rows, and 12 more variables: sched_arr_time
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <
## #   minute <dbl>, time_hour <dtm>
```

## Use desc() to sort by a variable in descending order

Sort by flights that departed with most delay:

In base R:

```
flights[order(flights$dep_delay, decreasing=TRUE),]
```

In dplyr:

```
arrange(flights, desc(dep_delay))
```

```
arrange(flights, desc(dep_delay))
```

```
## # A tibble: 336,776 x 19
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1  2013     1     9     641           900         1301    1242
## 2  2013     6    15    1432          1935         1137    1607
## 3  2013     1    10    1121          1635         1126    1239
## 4  2013     9    20    1139          1845         1014    1457
## 5  2013     7    22     845          1600         1005    1044
## 6  2013     4    10    1100          1900          960    1342
## 7  2013     3    17    2321           810          911     135
## 8  2013     6    27     959          1900          899    1236
## 9  2013     7    22    2257           759          898     121
## 10 2013    12     5     756          1700          896    1058
```

```
## # ... with 336,766 more rows, and 12 more variables: sched_arr_time
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <
## #   minute <dbl>, time_hour <dtm>
```

## Subsetting columns with `select()`

Keep only date and delay columns.

In base R:

```
flights[,c("year", "month", "day", "dep_delay", "arr_delay")]
```

In dplyr:

```
select(flights, year, month, day, dep_delay, arr_delay)
```

```
select(flights, year, month, day, dep_delay, arr_delay)
```

```
## # A tibble: 336,776 x 5
##   year month   day dep_delay arr_delay
##   <int> <int> <int>     <dbl>     <dbl>
## 1  2013     1     1         2         11
## 2  2013     1     1         4         20
## 3  2013     1     1         2         33
## 4  2013     1     1        -1        -18
## 5  2013     1     1        -6        -25
## 6  2013     1     1        -4         12
## 7  2013     1     1        -5         19
## 8  2013     1     1        -3        -14
## 9  2013     1     1        -3        -8
## 10 2013     1     1        -2         8
## # ... with 336,766 more rows
```

Keep first 9 columns (year through arr\_delay).

In base R:

```
flights[,1:9]
```

In dplyr:

```
select(flights, 1:9)
```

Keep first 9 columns (year through arr\_delay) by name.

In base R:

```
flights[,which(names(flights)== "year") :  
          which(names(flights)== "arr_delay")]
```

In dplyr:

```
select(flights, year:arr_delay)
```



```
select(flights, year:arr_delay)
```

```
## # A tibble: 336,776 x 9
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1  2013     1     1     517           515           2     830
## 2  2013     1     1     533           529           4     850
## 3  2013     1     1     542           540           2     923
## 4  2013     1     1     544           545          -1    1004
## 5  2013     1     1     554           600          -6     812
## 6  2013     1     1     554           558          -4     740
## 7  2013     1     1     555           600          -5     913
## 8  2013     1     1     557           600          -3     709
## 9  2013     1     1     557           600          -3     838
## 10 2013     1     1     558           600          -2     753
```

```
## # ... with 336,766 more rows, and 2 more variables: sched_arr_time <
```

```
## #   arr_delay <dbl>
```

Keep all columns except tail number and flight number.

In base R:

```
flights[, -c(which(names(flights)=="tailnum"),  
             which(names(flights)=="flight"))]
```

In dplyr:

```
select(flights, -tailnum, -flight)
```

```
select(flights, -tailnum, -flight)
```

```
## # A tibble: 336,776 x 17
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1  2013     1     1     517           515           2     830
## 2  2013     1     1     533           529           4     850
## 3  2013     1     1     542           540           2     923
## 4  2013     1     1     544           545          -1    1004
## 5  2013     1     1     554           600          -6     812
## 6  2013     1     1     554           558          -4     740
## 7  2013     1     1     555           600          -5     913
## 8  2013     1     1     557           600          -3     709
## 9  2013     1     1     557           600          -3     838
## 10 2013     1     1     558           600          -2     753
```

```
## # ... with 336,766 more rows, and 10 more variables: sched_arr_time
## #   arr_delay <dbl>, carrier <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>,
## #   time_hour <dtm>
```

## Useful functions for selecting column/variable names

- ▶ `dplyr::starts_with("arr")`
  - ▶ matches column names that begin with "arr"
- ▶ `dplyr::ends_with("time")`
  - ▶ matches column names that end with "time"
- ▶ `dplyr::contains("dep")`
  - ▶ matches column names that contain "dep"
- ▶ `dplyr::num_range("x", 1:3)`
  - ▶ matches "x1", "x2", and "x3"

Keep only columns starting with “arr”.

In base R:

```
flights[,substr(names(flights), 1, 3) == "arr"]
```

In dplyr:

```
select(flights, starts_with("arr"))
```

```
select(flights, starts_with("arr"))
```

```
## # A tibble: 336,776 x 2
##   arr_time arr_delay
##   <int>     <dbl>
## 1      830         11
## 2      850         20
## 3      923         33
## 4     1004        -18
## 5      812        -25
## 6      740         12
## 7      913         19
## 8      709        -14
## 9      838         -8
## 10     753          8
## # ... with 336,766 more rows
```

## Rename columns with `rename()`

`rename()` is a variant of `select()` that keeps all variables/columns while renaming the specified ones.

```
flights2 <- flights
names(flights2)[names(flights2)=="year" |
                 names(flights2)=="month" |
                 names(flights2)=="day"] <- c("YEAR", "MONTH", "DAY")
```

```
rename(flights, YEAR=year, MONTH=month, DAY=day)
```

```
rename(flights, YEAR=year, MONTH=month, DAY=day)
```

```
## # A tibble: 336,776 x 19
```

```
##   YEAR MONTH   DAY dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1  2013     1     1     517           515           2     830
## 2  2013     1     1     533           529           4     850
## 3  2013     1     1     542           540           2     923
## 4  2013     1     1     544           545          -1    1004
## 5  2013     1     1     554           600          -6     812
## 6  2013     1     1     554           558          -4     740
## 7  2013     1     1     555           600          -5     913
## 8  2013     1     1     557           600          -3     709
## 9  2013     1     1     557           600          -3     838
## 10 2013     1     1     558           600          -2     753
```

```
## # ... with 336,766 more rows, and 12 more variables: sched_arr_time
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <
## #   minute <dbl>, time_hour <dtm>
```



## New variables with mutate()

Create a new variable giving the average air speed (in mph) of each flight.

In base R:

```
flights2 <- flights  
flights2$speed <- flights2$distance / flights2$air_time * 60
```

In dplyr:

```
mutate(flights, speed = distance / air_time * 60)
```

```
mutate(flights, speed = distance / air_time * 60)
```

```
## # A tibble: 336,776 x 20
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>
## 1  2013     1     1     517           515           2     830
## 2  2013     1     1     533           529           4     850
## 3  2013     1     1     542           540           2     923
## 4  2013     1     1     544           545          -1    1004
## 5  2013     1     1     554           600          -6     812
## 6  2013     1     1     554           558          -4     740
## 7  2013     1     1     555           600          -5     913
## 8  2013     1     1     557           600          -3     709
## 9  2013     1     1     557           600          -3     838
## 10 2013     1     1     558           600          -2     753
```

```
## # ... with 336,766 more rows, and 13 more variables: sched_arr_time
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <
## #   minute <dbl>, time_hour <dtm>, speed <dbl>
```

Create two new variables giving (1) the average air speed (in mph) of each flight and (2) the amount of time gained in the air.

In base R:

```
flights2 <- flights  
flights2$speed <- flights2$distance / flights2$air_time * 60  
flights2$gain <- flights2$arr_delay - flights2$dep_delay
```

In dplyr:

```
mutate(flights,  
       speed = distance / air_time * 60,  
       gain = arr_delay - dep_delay)
```

```
mutate(flights,  
  speed = distance / air_time * 60,  
  gain = arr_delay - dep_delay)
```

```
## # A tibble: 336,776 x 21
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time  
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>  
## 1  2013     1     1     517           515           2     830  
## 2  2013     1     1     533           529           4     850  
## 3  2013     1     1     542           540           2     923  
## 4  2013     1     1     544           545          -1    1004  
## 5  2013     1     1     554           600          -6     812  
## 6  2013     1     1     554           558          -4     740  
## 7  2013     1     1     555           600          -5     913  
## 8  2013     1     1     557           600          -3     709  
## 9  2013     1     1     557           600          -3     838  
## 10 2013     1     1     558           600          -2     753
```

```
## # ... with 336,766 more rows, and 14 more variables: sched_arr_time  
## #   arr_delay <dbl>, carrier <chr>, flight <int>, tailnum <chr>,  
## #   origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>, hour <  
## #   minute <dbl>, time_hour <dtm>, speed <dbl>, gain <dbl>
```

# Anatomy of mutate()

```
mutate(data, var1 = expr1, var2 = expr2, ...)
```

- ▶ The first argument is the data
- ▶ The following arguments are named vectorized expressions that output a vector of the same length
- ▶ When used with `data.frames`, you can use variables created in the same `mutate()` call in the subsequent expressions

# Useful functions for creating variables

- ▶ Arithmetic such as `+`, `-`, `*`, `/`, `^`, etc.
  - ▶ These are vectorized and will recycle shorter variables
- ▶ Modular arithmetic such as `%%` and `%/%`
  - ▶ Useful for breaking apart integers (e.g., time into hours + minutes)
- ▶ Logs such as `log()`, `log2()` and `log10`
  - ▶ Useful for data with multiplicative variance
- ▶ Offsets such as `dplyr::lead()` and `dplyr::lag()`
  - ▶ Useful for running differences and data over time points
- ▶ Cumulative summaries such as `cumsum()`, `cumprod()`, `cummax()`, and `cummin()`
  - ▶ Also `dplyr::cummean()` for running means
- ▶ Logical operators such as `'=='`, `'!='`, `'>'`, `'<'`, `'>='`, `'<='`, etc.
  - ▶ Useful for turning continuous variables into categorical
- ▶ `dplyr::n()` gives the number of observations
  - ▶ `n()` can only be used inside `mutate()`, `filter()` and `summarise()`

## New variables with transmute()

`transmute()` is a variant of `mutate()` that keeps only the new variables and drops the rest.

In base R:

```
data.frame(distance = flights$distance,  
            speed = flights$distance / flights$air_time * 60,  
            gain = flights$arr_delay - flights$dep_delay,  
            gain_per_mile = (flights$arr_delay - flights$dep_delay) /  
                             flights$distance)
```

In dplyr:

```
transmute(flights,  
          distance = distance,  
          speed = distance / air_time * 60,  
          gain = arr_delay - dep_delay,  
          gain_per_mile = gain / distance)
```

```
transmute(flights,  
          distance = distance,  
          speed = distance / air_time * 60,  
          gain = arr_delay - dep_delay,  
          gain_per_mile = gain / distance)
```

```
## # A tibble: 336,776 x 4
```

```
##   distance speed  gain gain_per_mile  
##   <dbl> <dbl> <dbl>      <dbl>  
## 1    1400   370.     9      0.00643  
## 2    1416   374.    16      0.0113  
## 3    1089   408.    31      0.0285  
## 4    1576   517.   -17     -0.0108  
## 5     762   394.   -19     -0.0249  
## 6     719   288.    16      0.0223  
## 7    1065   404.    24      0.0225  
## 8     229   259.   -11     -0.0480  
## 9     944   405.    -5     -0.00530  
## 10     733   319.    10      0.0136  
## # ... with 336,766 more rows
```



## Summary statistics with summarise()

Get the mean departure and arrival delay.

In base R:

```
data.frame(mean_dep_delay = mean(flights$dep_delay,  
                                na.rm=TRUE),  
           mean_arr_delay = mean(flights$arr_delay,  
                                na.rm=TRUE))
```

In dplyr:

```
summarise(flights,  
          mean_dep_delay = mean(dep_delay, na.rm=TRUE),  
          mean_arr_delay = mean(arr_delay, na.rm=TRUE))
```

```
summarise(flights,  
          mean_dep_delay = mean(dep_delay, na.rm=TRUE),  
          mean_arr_delay = mean(arr_delay, na.rm=TRUE))
```

```
## # A tibble: 1 x 2  
##   mean_dep_delay mean_arr_delay  
##           <dbl>           <dbl>  
## 1           12.6             6.90
```

# Anatomy of summarise()

```
summarise(data, summary1 = expr1, summary2 = expr2, ...)
```

- ▶ The first argument is the data
- ▶ The following arguments are expressions that output a single value from a vector of values
- ▶ It is particularly important to consider missing values when summarizing data
- ▶ Also available as `summarize()`

# Useful functions for calculating summary statistics

- ▶ Measures of location such as `mean()` and `median()`
- ▶ Measures of spread such as `sd()`, `var()`, `IQR()`, and `mad()`
- ▶ Measures of rank such as `min()`, `max()`, and `quantile()`
- ▶ Counts such as:
  - ▶ `dplyr::n()` gives the number of observations
  - ▶ `sum(!is.na(x))` gives the number of non-missing values
  - ▶ `dplyr::n_distinct()` gives the number of unique values
- ▶ Remember that `sum(x == 10)` gives the count of `x == 10`
  - ▶ What does `mean(x == 10)` calculate?

Calculate the proportion of flights delayed more than 2 hours on arrival.

```
summarise(flights, mean(arr_delay > 120, na.rm=TRUE))
```

Calculate the number of unique airline carriers.

```
summarise(flights, n_distinct(carrier))
```

Calculate the proportion of flights with missing air times.

```
summarise(flights, sum(is.na(air_time)) / n())
```

```
summarise(flights, mean(arr_delay > 120, na.rm=TRUE))
```

```
## # A tibble: 1 x 1
##   `mean(arr_delay > 120, na.rm = TRUE)`
##                                     <dbl>
## 1                                     0.0307
```

```
summarise(flights, n_distinct(carrier))
```

```
## # A tibble: 1 x 1
##   `n_distinct(carrier)`
##               <int>
## 1                   16
```

```
summarise(flights, sum(is.na(air_time)) / n())
```

```
## # A tibble: 1 x 1
##   `sum(is.na(air_time))/n()`
##               <dbl>
## 1                   0.0280
```

## Grouped transformations with `group_by()`

`summarise()` and the other data manipulation verbs in `dplyr` become much more powerful when paired with `group_by()`.

Count the number of flights from each carrier.

```
summarise(group_by(flights, carrier), n())
```

Calculate the average arrival delay for each carrier.

```
summarise(group_by(flights, carrier),  
           mean(arr_delay, na.rm=TRUE))
```

You can group by multiple variables.

Use `ungroup()` to ungroup a grouped dataset.

```
summarise(group_by(flights, carrier),  
           mean(arr_delay, na.rm=TRUE))
```

```
## # A tibble: 16 x 2  
##   carrier `mean(arr_delay, na.rm = TRUE)`  
##   <chr>           <dbl>  
## 1 9E              7.38  
## 2 AA             0.364  
## 3 AS            -9.93  
## 4 B6             9.46  
## 5 DL             1.64  
## 6 EV            15.8  
## 7 F9            21.9  
## 8 FL            20.1  
## 9 HA            -6.92  
## 10 MQ            10.8  
## 11 OO            11.9  
## 12 UA             3.56  
## 13 US             2.13  
## 14 VX             1.76  
## 15 WN             9.65  
## 16 YV            15.6
```



## Piping with the pipe operator

Combining multiple `dplyr` verbs becomes much more expressive when used with the pipe operator `%>%`.

The pipe operator takes the return value of the expression on the LHS and turns it into the first argument of the function on the RHS.

```
foo(bar(baz(x)))
```

is the same as

```
baz(x) %>% bar() %>% foo()
```

is the same as

```
x %>% baz() %>% bar() %>% foo()
```

## Piping with the pipe operator (cont'd)

```
summarise(group_by(flights, carrier),  
           mean(arr_delay, na.rm=TRUE))
```

becomes

```
group_by(flights, carrier) %>%  
  summarise(mean(arr_delay, na.rm=TRUE))
```

or

```
flights %>%  
  group_by(carrier) %>%  
  summarise(mean(arr_delay, na.rm=TRUE))
```

```
flights %>%  
  group_by(carrier) %>%  
  summarise(mean(arr_delay, na.rm=TRUE))
```

```
## # A tibble: 16 x 2  
##   carrier `mean(arr_delay, na.rm = TRUE)`  
##   <chr>           <dbl>  
## 1 9E              7.38  
## 2 AA             0.364  
## 3 AS            -9.93  
## 4 B6             9.46  
## 5 DL             1.64  
## 6 EV            15.8  
## 7 F9            21.9  
## 8 FL            20.1  
## 9 HA            -6.92  
## 10 MQ            10.8  
## 11 OO            11.9  
## 12 UA             3.56  
## 13 US             2.13  
## 14 VX             1.76  
## 15 WN             9.65  
## 16 YV            15.6
```


# Data transformation + visualization

You can chain together dplyr verbs with ggplot2 too.

Visualize the average delay for each carrier.

```
library(ggplot2)
flights %>%
  group_by(carrier) %>%
  summarise(delay=mean(arr_delay, na.rm=TRUE)) %>%
  ggplot(aes(x=carrier, y=delay)) + geom_col()
```


```
flights %>%  
  group_by(carrier) %>%  
  summarise(delay=mean(arr_delay, na.rm=TRUE)) %>%  
  ggplot(aes(x=carrier, y=delay)) + geom_col()
```



For each destination, visualize the proportion of delayed arriving flights versus the average flight time in the air.

```
flights %>%  
  group_by(dest) %>%  
  summarise(prop_delayed = mean(arr_delay > 0, na.rm=TRUE),  
            flight_time = mean(air_time, na.rm=TRUE),  
            count = n()) %>%  
  ggplot(aes(x=flight_time, y=prop_delayed)) + geom_point()
```

```
## Warning: Removed 1 rows containing missing values (geom_point)
```



Map the number of flights from each destination to an aesthetic.

Add a smooth fitted line.


```
flights %>%  
  group_by(dest) %>%  
  summarise(prop_delayed = mean(arr_delay > 0, na.rm=TRUE),  
            flight_time = mean(air_time, na.rm=TRUE),  
            count = n()) %>%  
  ggplot(aes(x=flight_time,  
             y=prop_delayed)) +  
  geom_point(aes(size=count), alpha=1/2) +  
  geom_smooth()
```



```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```

```
## Warning: Removed 1 rows containing non-finite values (stat_smooth)
```

```
## Warning: Removed 1 rows containing missing values (geom_point)
```



# Exercises

- ▶ For carriers that flew more than 1000 flights in 2013, find the number of flights that weren't delayed on arrival.
- ▶ Find the average distance flown by each carrier in each month of 2013.
- ▶ Plot total miles flown each month in 2013.
- ▶ Plot the proportion of flights delayed by 10 minutes or more for each hour of the day.
- ▶ Plot the total distance flown versus the total time in arrival delays for each plane.
- ▶ Plot the relationship between the total time in the air and the total distance flown for each plane.
- ▶ Plot the average speed flown versus the average distance flown for each destination.
- ▶ Find the fastest plane.