

# R and SQL

Kylie Ariel Bemis

2/4/2019

# R and databases

Often, relational data is stored in a database managed by an RDBMS, so we need to know how to work with data in a database.

A database can also be a good solution for data which you do not want to load all into memory at once.

R and `dplyr` can interface with databases through the DBI package and a suitable backend:

- ▶ `RSQLite` works with SQLite databases
- ▶ `RMySQL` works with MySQL databases
- ▶ `RPostgreSQL` works with PostgreSQL databases

etc.

We can work with these with `dplyr` using the `dbplyr` package.

# R and databases with dbplyr

Suppose we wish to work with an SQLite database.

First we need to install the necessary packages.

```
install.packages(c("DBI", "RSQLite", "dbplyr"))
```

```
library(dbplyr)
```

```
##
```

```
## Attaching package: 'dbplyr'
```

```
## The following objects are masked from 'package:dplyr':
```

```
##
```

```
##      ident, sql
```

```
library(RSQLite)
```

## Connecting to a database in R

The DBLP is a database containing bibliographic data on major computer science journals and proceedings. A subset of it has been loaded into the SQLite database file “dblp.db”.

First, we need to open a *connection* to the database using `dbConnect()`.

The first argument of `dbConnect()` is the backend to use (provided by the RSQLite package in this case), and the second is the filepath to the database.

```
dbpath <- paste0("~/Documents/Northeastern/",  
                 "Courses/DS5110-Spring2019/data/",  
                 "DBLP-CSR-sqlite/dblp.db")  
con <- dbConnect(SQLite(), dbname=dbpath)
```

# Creating a tbl of the connected database

A `tbl` is the tidyverse's generalized notion of tabular data.

A tibble is a type of `tbl`. We can also create a `tbl` from a database.

```
dblp_main <- tbl(con, "general")  
dblp_authors <- tbl(con, "authors")
```

The first argument is the data source (our database connection).

The second argument is the name of the table within the database.

Our SQLite database contains the table “general”, which has information about papers published in computer science journals and proceedings.

```
dblp_main
```

```
## # Source:   table<general> [?? x 10]
```

```
## # Database: sqlite 3.22.0
```

```
## #   [/Users/kuwisdeltu/Dropbox/Northeastern/Courses/DS5110-Spr
```

```
##   k          year conf  crossref      cs      de      se      th publi
```

```
##   <chr>      <int> <chr> <chr>      <int> <int> <int> <int> <chr>
```

```
##  1 conf/a~  2013 AAAI  conf/aa~      1      0      0      0 AAAI
```

```
##  2 conf/a~  2015 AAAI  conf/aa~      1      0      0      0 AAAI
```

```
##  3 conf/a~  2014 AAAI  conf/aa~      1      0      0      0 AAAI
```

```
##  4 conf/a~  2015 AAAI  conf/aa~      1      0      0      0 AAAI
```

```
##  5 conf/a~  2015 AAAI  conf/aa~      1      0      0      0 AAAI
```

```
##  6 conf/a~  2014 AAAI  conf/aa~      1      0      0      0 AAAI
```

```
##  7 conf/a~  2015 AAAI  conf/aa~      1      0      0      0 AAAI
```

```
##  8 conf/a~  2015 AAAI  conf/aa~      1      0      0      0 AAAI
```

```
##  9 conf/a~  2014 AAAI  conf/aa~      1      0      0      0 AAAI
```

```
## 10 conf/a~  2015 AAAI  conf/aa~      1      0      0      0 AAAI
```

```
## # ... with more rows
```

```
dblp_authors
```

```
## # Source:   table<authors> [?? x 6]
```

```
## # Database: sqlite 3.22.0
```

```
## #    [/Users/kuwisdeltu/Dropbox/Northeastern/Courses/DS5110-Spr
```

```
##      id k                                pos name                                gende
```

```
##      <int> <chr>                        <int> <chr>                        <chr>
```

```
## 1         1 conf/aaai/0001M13             0 Chang Wang 0001      M
```

```
## 2         2 conf/aaai/0001M13             1 Sridhar Mahadevan    M
```

```
## 3         3 conf/aaai/0001T15             0 Claudia Schulz 0001  F
```

```
## 4         4 conf/aaai/0001T15             1 Francesca Toni      F
```

```
## 5         5 conf/aaai/0001TZLL14          0 Jing Zhang 0001     F
```

```
## 6         6 conf/aaai/0001TZLL14          1 Jie Tang            F
```

```
## 7         7 conf/aaai/0001TZLL14          2 Honglei Zhuang      -
```

```
## 8         8 conf/aaai/0001TZLL14          3 Cane Wing-ki Leung  F
```

```
## 9         9 conf/aaai/0001TZLL14          4 Juan-Zi Li          -
```

```
## 10        10 conf/aaai/0001VD15           0 Bart Bogaerts 0001  M
```

```
## # ... with more rows
```

We can perform dplyr operations on the database using dplyr verbs.

```
dblp_main %>% summarise(cs_papers=sum(cs))
```

```
## Warning: Missing values are always removed in SQL.  
## Use `SUM(x, na.rm = TRUE)` to silence this warning
```

```
## # Source:   lazy query [?? x 1]  
## # Database: sqlite 3.22.0  
## #   [/Users/kuwisdellu/Dropbox/Northeastern/Courses/DS5110-Spr  
##   cs_papers  
##           <int>  
## 1         96823
```



```
dblp_main %>%  
  filter(year > 2010) %>%  
  group_by(year) %>%  
  summarise(cs_papers=sum(cs))
```

```
## Warning: Missing values are always removed in SQL.  
## Use `SUM(x, na.rm = TRUE)` to silence this warning
```

```
## # Source:   lazy query [?? x 2]  
## # Database: sqlite 3.22.0  
## #   [/Users/kuwisdelu/Dropbox/Northeastern/Courses/DS5110-Spr  
##   year cs_papers  
##   <int>      <int>  
## 1  2011      6060  
## 2  2012      5772  
## 3  2013      7391  
## 4  2014      6538  
## 5  2015      5910
```

When working with databases, dbplyr tries to offload as much work as possible to the database itself, and delay execution until we need the result.

We can see the actual SQL commands generated using `show_query()`

```
query <- dblp_main %>%  
  filter(year > 2010) %>%  
  group_by(year) %>%  
  summarise(cs_papers=sum(cs))  
show_query(query)
```

```
## Warning: Missing values are always removed in SQL.  
## Use `SUM(x, na.rm = TRUE)` to silence this warning
```

```
## <SQL>  
## SELECT `year`, SUM(`cs`) AS `cs_papers`  
## FROM `general`  
## WHERE (`year` > 2010.0)  
## GROUP BY `year`
```

## Collecting a query into R

The query is not executed until the result is needed. You can use `collect()` to execute the query and pull the result into R.

```
query %>% collect()
```

```
## Warning: Missing values are always removed in SQL.  
## Use `SUM(x, na.rm = TRUE)` to silence this warning
```

```
## # A tibble: 5 x 2  
##   year cs_papers  
##   <int>   <int>  
## 1  2011     6060  
## 2  2012     5772  
## 3  2013     7391  
## 4  2014     6538  
## 5  2015     5910
```

## Pulling database data into R

Not all databases support all dplyr data manipulation verbs, so if necessary (and the data is small enough), we can always pull the data into R as a tibble and work that way.

```
dblp_main %>% collect()
```

```
## # A tibble: 148,521 x 10
```

```
##   k          year conf crossref      cs      de      se      th publi
```

```
##   <chr>    <int> <chr> <chr>    <int> <int> <int> <int> <chr>
```

```
## 1 conf/a~  2013 AAAI  conf/aa~    1      0      0      0 AAAI
```

```
## 2 conf/a~  2015 AAAI  conf/aa~    1      0      0      0 AAAI
```

```
## 3 conf/a~  2014 AAAI  conf/aa~    1      0      0      0 AAAI
```

```
## 4 conf/a~  2015 AAAI  conf/aa~    1      0      0      0 AAAI
```

```
## 5 conf/a~  2015 AAAI  conf/aa~    1      0      0      0 AAAI
```

```
## 6 conf/a~  2014 AAAI  conf/aa~    1      0      0      0 AAAI
```

```
## 7 conf/a~  2015 AAAI  conf/aa~    1      0      0      0 AAAI
```

```
## 8 conf/a~  2015 AAAI  conf/aa~    1      0      0      0 AAAI
```

```
## 9 conf/a~  2014 AAAI  conf/aa~    1      0      0      0 AAAI
```

```
## 10 conf/a~ 2015 AAAI  conf/aa~    1      0      0      0 AAAI
```

```
## # ... with 148,511 more rows
```

## SQL equivalents to dplyr functions

Most functions in dplyr are simple and flexible versions of statements in SQL.

- ▶ `select()` and `SELECT` are used to select columns from a table
- ▶ `filter()` and `WHERE` are used to subset a table by on rows based on given conditions
- ▶ `arrange()` and `ORDER BY` are used to order the rows of a table
- ▶ `mutate()` does not have a direct equivalent, but `AS` performs a similar function of returning columns (potentially transformed) with a new name
- ▶ `summarise()` does not have a direct equivalent, but SQL provides several summary functions such as `SUM`, `AVG`, `COUNT`, etc.

Let's see some ways to perform the same operations using dplyr and SQL.

## select() and SELECT

```
dblp_authors %>% select(k, name, gender)
```

```
query <- con %>%  
  dbSendQuery("SELECT k, name, gender  
              FROM authors")  
result <- as_tibble(dbFetch(query))  
dbClearResult(query)  
result
```

```
## # Source: lazy query [?? x 3]
## # Database: sqlite 3.22.0
## # [/Users/kuwisdeltu/Dropbox/Northeastern/Courses/DS5110-Spr
## k name gender
## <chr> <chr> <chr>
## 1 conf/aaai/0001M13 Chang Wang 0001 M
## 2 conf/aaai/0001M13 Sridhar Mahadevan M
## 3 conf/aaai/0001T15 Claudia Schulz 0001 F
## 4 conf/aaai/0001T15 Francesca Toni F
## 5 conf/aaai/0001TZLL14 Jing Zhang 0001 F
## 6 conf/aaai/0001TZLL14 Jie Tang F
## 7 conf/aaai/0001TZLL14 Honglei Zhuang -
## 8 conf/aaai/0001TZLL14 Cane Wing-ki Leung F
## 9 conf/aaai/0001TZLL14 Juan-Zi Li -
## 10 conf/aaai/0001VD15 Bart Bogaerts 0001 M
## # ... with more rows
```

## filter() and WHERE

```
dblp_authors %>% filter(gender == "F")
```

```
query <- con %>%  
  dbSendQuery("SELECT *  
              FROM authors  
              WHERE gender = 'F'")  
result <- as_tibble(dbFetch(query))  
dbClearResult(query)  
result
```



## # Source: lazy query [?? x 6]

## # Database: sqlite 3.22.0

## # [/Users/kuwisdeltu/Dropbox/Northeastern/Courses/DS5110-Spr

##	id	k	pos	name	gend
----	----	---	-----	------	------

##	<int>	<chr>	<int>	<chr>	<chr>
----	-------	-------	-------	-------	-------

##	1	3	0	conf/aaai/0001T15	Claudia Schulz 0001 F
----	---	---	---	-------------------	-----------------------

##	2	4	1	conf/aaai/0001T15	Francesca Toni F
----	---	---	---	-------------------	------------------

##	3	5	0	conf/aaai/0001TZLL14	Jing Zhang 0001 F
----	---	---	---	----------------------	-------------------

##	4	6	1	conf/aaai/0001TZLL14	Jie Tang F
----	---	---	---	----------------------	------------

##	5	8	3	conf/aaai/0001TZLL14	Cane Wing-ki Leung F
----	---	---	---	----------------------	----------------------

##	6	17	1	conf/aaai/0002GYSZL14	Hua Guo F
----	---	----	---	-----------------------	-----------

##	7	18	2	conf/aaai/0002GYSZL14	Yi Yang F
----	---	----	---	-----------------------	-----------

##	8	27	0	conf/aaai/0003MGF14	Wei Li 0002 F
----	---	----	---	---------------------	---------------

##	9	29	2	conf/aaai/0003MGF14	Tovi Grossman F
----	---	----	---	---------------------	-----------------

##	10	33	2	conf/aaai/0005YJZ15	Rong Jin F
----	----	----	---	---------------------	------------

## # ... with more rows

## filter() and WHERE (cont'd)

```
dblp_authors %>% filter(gender == "F" & prob > 0.95)
```

```
query <- con %>%  
  dbSendQuery("SELECT *  
               FROM authors  
               WHERE gender = 'F' AND prob > 0.95")  
result <- as_tibble(dbFetch(query))  
dbClearResult(query)  
result
```

## # Source: lazy query [?? x 6]

## # Database: sqlite 3.22.0

## # [/Users/kuwisdeldu/Dropbox/Northeastern/Courses/DS5110-Spr

##	id	k	pos	name	g
----	----	---	-----	------	---

##	<int>	<chr>	<int>	<chr>	<
----	-------	-------	-------	-------	---

##	1	3	conf/aaai/0001T15	0	Claudia Schulz 0001 F
----	---	---	-------------------	---	-----------------------

##	2	4	conf/aaai/0001T15	1	Francesca Toni F
----	---	---	-------------------	---	------------------

##	3	29	conf/aaai/0003MGF14	2	Tovi Grossman F
----	---	----	---------------------	---	-----------------

##	4	39	conf/aaai/Abbott88	0	Kathy H. Abbott F
----	---	----	--------------------	---	-------------------

##	5	52	conf/aaai/AbellaK93	0	Alicia Abella F
----	---	----	---------------------	---	-----------------

##	6	54	conf/aaai/AbergALS06	0	Cécile Aberg F
----	---	----	----------------------	---	----------------

##	7	58	conf/aaai/AbernethySB08	0	Jennifer Abernethy F
----	---	----	-------------------------	---	----------------------

##	8	60	conf/aaai/AbernethySB08	2	Elizabeth Bradley F
----	---	----	-------------------------	---	---------------------

##	9	66	conf/aaai/AbtahiF11	0	Farnaz Abtahi F
----	---	----	---------------------	---	-----------------

##	10	70	conf/aaai/Abu-HakimaHP91	2	Sieu Phan F
----	----	----	--------------------------	---	-------------

## # ... with more rows

## arrange() and ORDER BY

```
dblp_authors %>% arrange(name)
```

```
query <- con %>%  
  dbSendQuery("SELECT *  
              FROM authors  
              ORDER BY name")  
result <- as_tibble(dbFetch(query))  
dbClearResult(query)  
result
```

```
## # Source:      table<authors> [?? x 6]
## # Database:    sqlite 3.22.0
## #      [/Users/kuwisdeltu/Dropbox/Northeastern/Courses/DS5110-Spr
## # Ordered by: name
```

##	id	k	pos	name	gender
##	<int>	<chr>	<int>	<chr>	<chr>
##	1	134173	conf/dolap/Mangisengi	T98	1 A Min Tjoa -
##	2	134238	conf/dolap/Nguyen	ST05	2 A Min Tjoa -
##	3	140374	conf/er/Eder	K86	2 A Min Tjoa -
##	4	142860	conf/er/Pernul	T91	1 A Min Tjoa -
##	5	142863	conf/er/Pernul	WT93	2 A Min Tjoa -
##	6	143978	conf/er/Tjoa	B93	0 A Min Tjoa -
##	7	143980	conf/er/Tjoa	W79	0 A Min Tjoa -
##	8	181589	conf/icde/Esmayr	KPPT96	4 A Min Tjoa -
##	9	187808	conf/icde/Schrefl	TW84	1 A Min Tjoa -
##	10	9495	conf/aaai/Kuan	PH86	2 A-Chuan Hsueh -

```
## # ... with more rows
```

## mutate() and AS

```
dblp_authors %>%  
  mutate(pos1 = pos + 1) %>%  
  select(k, pos1, name)
```

```
query <- con %>%  
  dbSendQuery("SELECT k, pos + 1 AS pos1, name  
              FROM authors")  
result <- as_tibble(dbFetch(query))  
dbClearResult(query)  
result
```

```
## # Source: lazy query [?? x 3]
## # Database: sqlite 3.22.0
## # [/Users/kuwisdeltu/Dropbox/Northeastern/Courses/DS5110-Spr
## k pos1 name
## <chr> <dbl> <chr>
## 1 conf/aaai/0001M13 1 Chang Wang 0001
## 2 conf/aaai/0001M13 2 Sridhar Mahadevan
## 3 conf/aaai/0001T15 1 Claudia Schulz 0001
## 4 conf/aaai/0001T15 2 Francesca Toni
## 5 conf/aaai/0001TZLL14 1 Jing Zhang 0001
## 6 conf/aaai/0001TZLL14 2 Jie Tang
## 7 conf/aaai/0001TZLL14 3 Honglei Zhuang
## 8 conf/aaai/0001TZLL14 4 Cane Wing-ki Leung
## 9 conf/aaai/0001TZLL14 5 Juan-Zi Li
## 10 conf/aaai/0001VD15 1 Bart Bogaerts 0001
## # ... with more rows
```

summarise() and COUNT, AVG, SUM, etc.

```
dblp_authors %>% summarize(mean(prob))
```

```
query <- con %>%  
  dbSendQuery("SELECT AVG(prob)  
              FROM authors")  
result <- as_tibble(dbFetch(query))  
dbClearResult(query)  
result
```



```
## Warning: Missing values are always removed in SQL.  
## Use `AVG(x, na.rm = TRUE)` to silence this warning
```

```
## # Source:   lazy query [?? x 1]  
## # Database: sqlite 3.22.0  
## #    [/Users/kuwisdeltu/Dropbox/Northeastern/Courses/DS5110-Spr  
##   `mean(prob)`  
##           <dbl>  
## 1           1.12
```

group\_by() and COUNT, AVG, SUM, etc.

```
dblp_authors %>%  
  group_by(gender) %>%  
  summarize(n())
```

```
query <- con %>%  
  dbSendQuery("SELECT gender, COUNT() AS n  
              FROM authors  
              GROUP BY gender")  
result <- as_tibble(dbFetch(query))  
dbClearResult(query)  
result
```

```
## # Source: lazy query [?? x 2]
## # Database: sqlite 3.22.0
## # [/Users/kuwisdelu/Dropbox/Northeastern/Courses/DS5110-Spr
## gender `n()`
## <chr> <int>
## 1 - 59009
## 2 F 69500
## 3 M 287936
```

# SQL JOINS

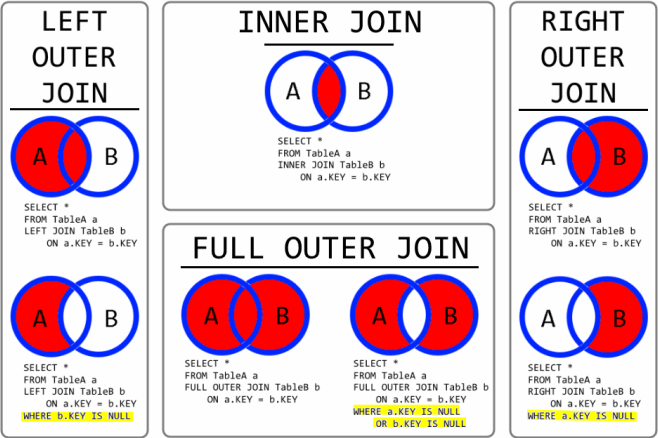


Figure 1:

```
x <- tribble(  
  ~key, ~val_x,  
    1, "x1",  
    2, "x2",  
    3, "x3"  
)  
y <- tribble(  
  ~key, ~val_y,  
    1, "y1",  
    2, "y2",  
    4, "y3"  
)
```

```
con2 <- dbConnect(SQLite(), ":memory:")
dbWriteTable(con2, "x", x)
dbWriteTable(con2, "y", y)
x <- tbl(con2, "x")
y <- tbl(con2, "y")
```

## Inner joins

```
inner_join(x, y, by="key")

query <- con2 %>%
  dbSendQuery("SELECT x.key AS key, val_x, val_y
              FROM x
              INNER JOIN y USING(key)")
result <- as_tibble(dbFetch(query))
dbClearResult(query)
result
```

```
inner_join(x, y)
```

```
## Joining, by = "key"
```

```
## # Source:   lazy query [?? x 3]
```

```
## # Database: sqlite 3.22.0 [:memory:]
```

```
##      key val_x val_y
```

```
##    <dbl> <chr> <chr>
```

```
## 1      1 x1     y1
```

```
## 2      2 x2     y2
```



## Inner joins (cont'd)

```
inner_join(x, y, by=c("key", "key"))
```

```
query <- con2 %>%  
  dbSendQuery("SELECT x.key AS key, val_x, val_y  
               FROM x  
               INNER JOIN y ON x.key = y.key")  
result <- as_tibble(dbFetch(query))  
dbClearResult(query)  
result
```

```
inner_join(x, y, by=c("key", "key"))
```

```
## # Source:   lazy query [?? x 3]
## # Database: sqlite 3.22.0 [:memory:]
##      key val_x val_y
##    <dbl> <chr> <chr>
## 1      1 x1     y1
## 2      2 x2     y2
```

## Left joins

```
left_join(x, y)
```

```
query <- con2 %>%  
  dbSendQuery("SELECT x.key AS key, val_x, val_y  
                FROM x  
                LEFT JOIN y ON x.key = y.key")  
result <- as_tibble(dbFetch(query))  
dbClearResult(query)  
result
```

```
left_join(x, y)
```

```
## Joining, by = "key"
```

```
## # Source:   lazy query [?? x 3]
```

```
## # Database: sqlite 3.22.0 [:memory:]
```

```
##      key val_x val_y
```

```
##    <dbl> <chr> <chr>
```

```
## 1      1 x1     y1
```

```
## 2      2 x2     y2
```

```
## 3      3 x3     <NA>
```

## Left excluding joins

```
x %>% left_join(y) %>% anti_join(y)
```

```
query <- con2 %>%  
  dbSendQuery("SELECT x.key AS key, val_x, val_y  
               FROM x  
               LEFT JOIN y ON x.key = y.key  
               WHERE y.key IS NULL")  
result <- as_tibble(dbFetch(query))  
dbClearResult(query)  
result
```

```
x %>% left_join(y) %>% anti_join(y)
```

```
## Joining, by = "key"
```

```
## Joining, by = c("key", "val_y")
```

```
## # Source:   lazy query [?? x 3]
```

```
## # Database: sqlite 3.22.0 [:memory:]
```

```
##      key val_x val_y
```

```
##    <dbl> <chr> <chr>
```

```
## 1      3 x3    <NA>
```

## Right joins

*# not supported by SQLite*

```
right_join(x, y)
```

*# not supported by SQLite*

```
query <- con2 %>%
```

```
  dbSendQuery("SELECT x.key AS key, val_x, val_y  
               FROM x  
               RIGHT JOIN y ON x.key = y.key")
```

```
right_join(x, y)
```

```
## Joining, by = "key"
```

```
## Error in result_create(conn@ptr, statement): RIGHT and FULL O
```



# Thoughts on SQL vs dplyr

Some basic SQL rules:

- ▶ SQL statements that retrieve data from the database begin with `SELECT`; use `*` to select all columns
  - ▶ You don't need to use `dplyr::select()` unless you're subsetting columns
- ▶ In SQL statements, you always need to specify the table(s) you are manipulating
  - ▶ In `dplyr`, each table will usually be a separate variable
- ▶ SQL expects clauses following `SELECT` to be in a particular order
  - ▶ You can apply `dplyr` functions in any order (though results may differ)

# Thoughts on SQL vs dplyr (cont'd)

Advantages and disadvantages:

- ▶ It is easy to chain functions separately in dplyr and see intermediate results; this may be more difficult in SQL
- ▶ Reasoning about tables and operations may be easier when you can see intermediate results like this
- ▶ Some complex SQL operations are not supported by dplyr directly
- ▶ If the data fits in memory, you can use R programming to perform more complex operations
- ▶ If the data does not fit in memory, you may need to use SQL to perform the complex operation on the database
- ▶ Use whichever is easier for you, then pull data into R for visualization and analysis.

It is always good to be familiar with all available tools. Most data science workflows will involve incorporating multiple tools

# Relational data examples with DBLP

Let's go back to the DBLP dataset for more practice with dbplyr.

Calculate the total number of distinct authors in the dataset.


```
dblp_authors %>%  
  summarise(num_papers = n(),  
            num_authors = n_distinct(name))
```

```
## # Source:   lazy query [?? x 2]  
## # Database: sqlite 3.22.0  
## #       [/Users/kuwisdeltu/Dropbox/Northeastern/Courses/DS5110-Spr  
##   num_papers num_authors  
##           <int>         <int>  
## 1       416445       126094
```

Calculate and plot the number of distinct authors published each year.

```
dblp_authors %>%  
  left_join(dblp_main) %>%  
  group_by(year) %>%  
  summarise(num_authors = n_distinct(name)) %>%  
  collect() %>%  
  ggplot() +  
  geom_col(aes(x=year, y=num_authors))
```

```
## Joining, by = "k"
```




Calculate and plot the number of papers published in CS, DE, SE, and TH.

```
dblp_tidy <- dblp_authors %>%  
  left_join(dblp_main) %>%  
  select(k, year, conf, name, gender, prob, cs, de, se, th) %>%  
  collect() %>%  
  gather(key="area", value="indicator", cs, de, se, th) %>%  
  filter(indicator == 1) %>%  
  select(-indicator) %>%  
  filter(gender %in% c('M', 'F'))
```


Some papers may belong to multiple areas, so we will count such a paper multiple times, once for each area. Since we are primarily interested in comparing between different areas of computer science, this is fine. However, if we wanted to count each paper only once, a different approach (adding an additional column for “interdisciplinary”?) would be required.

```
ggplot(dblp_tidy) + geom_bar(aes(x=area))
```



Plot the number of papers published in CS, DE, SE, and TH by gender.


```
ggplot(dblp_tidy) + geom_bar(aes(x=area, fill=gender))
```





Plot the number of distinct men and women published in CS, DE, SE, and TH.

```
dblp_tidy %>%  
  group_by(gender, area) %>%  
  summarise(count = n_distinct(name)) %>%  
  collect() %>%  
  ggplot() +  
  geom_col(aes(x=area, y=count, fill=gender))
```



Plot the proportion of distinct men and women published in CS, DE, SE, and TH.

```
auth_area <- dblp_tidy %>%  
  group_by(area) %>%  
  summarise(total = n_distinct(name))  
  
genauth_area <- dblp_tidy %>%  
  group_by(gender, area) %>%  
  summarise(gencount = n_distinct(name))  
  
left_join(genauth_area, auth_area) %>%  
  mutate(genprop = gencount / total) %>%  
  ggplot() +  
  geom_col(aes(x=area,  
               y=genprop,  
               fill=gender))
```

