

Introduction to Text Mining

Kylie Ariel Bemis

11/2/2018

Mining text data in R using tidytext




Figure 1: *Text Mining with R* by Julia Silge and David Robinson

Book freely available at <http://tidytextmining.com>.

Text data

Typically, text data is not tidy.

Often, text data is not even tabular or structured at all.

Most commonly, raw text data comes in two forms:

- ▶ **String:** Text data (e.g., emails, tweets, books) can be stored as strings and is most often first encountered and read into memory as strings, i.e., a character vector in R.
- ▶ **Corpus:** A collection of text data, typically stored as strings, where each object is distinct in some way (i.e., each string is a specific book or email or tweet), sometimes annotated with additional metadata

To analyze text data, we first need to put it into some sort of structured, tabular format.

Tidying text data

To create a structured dataset from unstructured text data, first we must identify the basic unit of text that we need to analyze.

A **token** (or term) is a meaningful unit of text that we are interested in using for analysis; we apply **tokenization** to split the text into tokens.

Most commonly, tokens are:

- ▶ A **word** is a distinct, meaningful element of text. In English, words are easy to identify, but in other languages, word tokenization can be a more difficult process.
- ▶ An **n-gram** is a contiguous sequence of words in a text. E.g., a 2-gram or “bigram” is two consecutive words, and a 3-gram or “trigram” is three consecutive words.

Tokenization example

```
library(tokenizers)
text <- "The quick brown fox jumps over the lazy dog"
tokenize_words(text)
```

```
## [[1]]
## [1] "the"    "quick" "brown" "fox"    "jumps" "over"  "the"
## [8] "lazy"   "dog"
```

```
tokenize_ngrams(text, n=2)
```

```
## [[1]]
## [1] "the quick"    "quick brown" "brown fox"    "fox jumps"
## [5] "jumps over"   "over the"     "the lazy"     "lazy dog"
```

Tidying text data (cont'd)

After a text document has been tokenized, there are two common formats for working with the text data:

- ▶ The **tidy text** format follows the idea of “tidy data” and formats text data as “one-token-per-row”. This is the format we will focus on now.
 - ▶ Each token is a row
 - ▶ Each variable is a column
 - ▶ Each value is in a cell
- ▶ A **document-term matrix** is a sparse matrix format used commonly in machine learning where each *unique* token (or term) is a column.
 - ▶ Each document is a row
 - ▶ Each term is a column
 - ▶ Each value is a count or frequency

These formats are each useful for different purposes. We will focus on the tidy format, because it is compatible with all of the tidyverse functions we have learned so far.

Tidy text format

```
emily <- c("Because I could not stop for Death -",  
          "He kindly stopped for me -",  
          "The Carriage held but just Ourselves -",  
          "and Immortality")
```

```
## # A tibble: 20 x 2  
##   line word  
##   <int> <chr>  
## 1     1 because  
## 2     1 i  
## 3     1 could  
## 4     1 not  
## 5     1 stop  
## 6     1 for  
## 7     1 death  
## 8     2 he  
## 9     2 kindly  
## 10    2 stopped  
## # ... with 10 more rows
```

Document-term matrix

```
emily <- c("Because I could not stop for Death -",  
           "He kindly stopped for me -",  
           "The Carriage held but just Ourselves -",  
           "because Immortality")
```

```
##   because i could not stop for death he kindly stopped me  
## 1      1 1      1 1      1 1      1 0      0      0 0  
## 2      0 0      0 0      0 1      0 1      1      1 1  
## 3      0 0      0 0      0 0      0 0      0      0 0  
## 4      1 0      0 0      0 0      0 0      0      0 0  
##   the carriage held but just ourselves immortality  
## 1      0      0 0 0 0      0      0  
## 2      0      0 0 0 0      0      0  
## 3      1      1 1 1 1      1      0  
## 4      0      0 0 0 0      0      1
```


Tokenizing with `unnest_tokens()`

Given a data frame with a column of strings, `unnest_tokens()` will tokenize the strings and return a new data frame in tidy text format.

```
unnest_tokens(data, output, input, token, ...)
```

- ▶ The first argument is the data
- ▶ The output is the name of a new column that will be created with the tokens
- ▶ The input is the name of the column with the strings to be tokenized
- ▶ The token is the type of tokenizer to use (e.g., word or n-gram)

unnest_tokens() example

```
library(tidytext)
poem <- tibble(line=1:4, text=emily)
tidy_poem <- unnest_tokens(poem, word, text)
print(tidy_poem, n=10)
```

```
## # A tibble: 20 x 2
##   line word
##   <int> <chr>
## 1     1 1 because
## 2     2 1 i
## 3     3 1 could
## 4     4 1 not
## 5     5 1 stop
## 6     6 1 for
## 7     7 1 death
## 8     8 2 he
## 9     9 2 kindly
## 10    10 2 stopped
## # ... with 10 more rows
```

unnest_tokens() example (cont'd)

```
library(tidytext)
poem <- tibble(line=1:4, text=emily)
tidy_poem <- unnest_tokens(poem, bigram, text, token="ngrams", n=2)
print(tidy_poem, n=10)
```

```
## # A tibble: 16 x 2
##   line bigram
##   <int> <chr>
## 1     1 1 because i
## 2     1 1 i could
## 3     1 1 could not
## 4     1 1 not stop
## 5     1 1 stop for
## 6     1 1 for death
## 7     2 2 he kindly
## 8     2 2 kindly stopped
## 9     2 2 stopped for
## 10    2 2 for me
## # ... with 6 more rows
```

Importing text data into R

In order to analyze text data, we must first import it into R as strings. We can do this simply with the `readLines()` function.

```
library(tidytext)
text <- readLines("prideprejudice.txt")
head(text, n=10)
```

```
## [1] "PRIDE AND PREJUDICE"
## [2] ""
## [3] "By Jane Austen"
## [4] ""
## [5] ""
## [6] ""
## [7] "Chapter 1"
## [8] ""
## [9] ""
## [10] "It is a truth universally acknowledged, that a single man in p
```

Preparing text data for tidying

In order to use the `unnest_tokens()` function, we need to put the string data in a data frame.

```
prideprejudice <- tibble(line=1:length(text), text=text)
prideprejudice
```

```
## # A tibble: 13,030 x 2
##   line text
##   <int> <chr>
## 1     1 PRIDE AND PREJUDICE
## 2     2 ""
## 3     3 By Jane Austen
## 4     4 ""
## 5     5 ""
## 6     6 ""
## 7     7 Chapter 1
## 8     8 ""
## 9     9 ""
## 10    10 It is a truth universally acknowledged, that a si~
## # ... with 13,020 more rows
```

Tidying the text data

```
tidy_pride <- prideprejudice %>%  
  unnest_tokens(word, text)
```

```
tidy_pride
```

```
## # A tibble: 122,204 x 2  
##       line word  
##   <int> <chr>  
## 1      1  pride  
## 2      1 and  
## 3      1 prejudice  
## 4      3 by  
## 5      3 jane  
## 6      3 austen  
## 7      7 chapter  
## 8      7 1  
## 9     10 it  
## 10     10 is  
## # ... with 122,194 more rows
```

Check most common words

```
tidy_pride %>% count(word, sort=TRUE)
```

```
## # A tibble: 6,538 x 2
##   word      n
##   <chr> <int>
## 1 the    4331
## 2 to     4162
## 3 of     3610
## 4 and    3585
## 5 her    2203
## 6 i      2065
## 7 a      1954
## 8 in     1880
## 9 was    1843
## 10 she   1695
## # ... with 6,528 more rows
```

Not very useful.

Remove stop words

Stop words are words that are not useful for analysis.

```
stop_words
```

```
## # A tibble: 1,149 x 2
##   word      lexicon
##   <chr>    <chr>
## 1 a       SMART
## 2 a's     SMART
## 3 able    SMART
## 4 about   SMART
## 5 above   SMART
## 6 according SMART
## 7 accordingly SMART
## 8 across  SMART
## 9 actually SMART
## 10 after  SMART
## # ... with 1,139 more rows
```


Check most common words again

```
pride_stop <- tibble(word=c("elizabeth", "darcy", "bennet", "jan  
  
tidy_pride %>%  
  anti_join(stop_words, by="word") %>%  
  anti_join(pride_stop, by="word") %>%  
  count(word, sort=TRUE)
```

```
## # A tibble: 6,013 x 2
```

```
##   word      n
```

```
##   <chr>   <int>
```

```
## 1 miss    283
```

```
## 2 time    203
```

```
## 3 lady    183
```

```
## 4 sister  180
```

```
## 5 wickham 162
```

```
## 6 dear    158
```

```
## 7 collins 156
```

```
## 8 family  151
```


```
## 9 day     140
```

```
## 10 lydia  133
```

Visualize the most common words

```
tidy_pride %>%  
  anti_join(stop_words, by="word") %>%  
  count(word, sort=TRUE) %>%  
  filter(n > 150) %>%  
  mutate(word = reorder(word, n)) %>%  
  ggplot(aes(x=word, y=n)) +  
  geom_col() +  
  coord_flip()
```

What does `reorder()` do here?



Analyzing all Jane Austen novels

Now that we know how to import and tidy raw text data, we will use data from the `janeaustenr` package, which includes the full text from all six of Jane Austen's completed, published novels.

```
library(janeaustenr)
austen_books()
```

```
## # A tibble: 73,422 x 2
##   text                                book
##   * <chr>                            <fct>
## 1 SENSE AND SENSIBILITY Sense & Sensibility
## 2 ""                               Sense & Sensibility
## 3 by Jane Austen              Sense & Sensibility
## 4 ""                               Sense & Sensibility
## 5 (1811)                       Sense & Sensibility
## 6 ""                               Sense & Sensibility
## 7 ""                               Sense & Sensibility
## 8 ""                               Sense & Sensibility
## 9 ""                               Sense & Sensibility
## 10 CHAPTER 1                  Sense & Sensibility
## # ... with 73,412 more rows
```

Tidying the works of Jane Austen

```
library(stringr)
tidy_austen <- austen_books() %>%
  group_by(book) %>%
  mutate(linenumber = row_number(),
         chapter = cumsum(str_detect(text,
                                     regex("^chapter [\\divxlc]",
                                           ignore_case = TRUE)))) %>%
  ungroup() %>%
  unnest_tokens(word, text)
```

```
tidy_austen
```


```
## # A tibble: 725,055 x 4
```

```
##   book                linenumber chapter word
##   <fct>                <int>    <int> <chr>
## 1 Sense & Sensibility      1        0 sense
## 2 Sense & Sensibility      1        0 and
## 3 Sense & Sensibility      1        0 sensibility
## 4 Sense & Sensibility      3        0 by
## 5 Sense & Sensibility      3        0 jane
## 6 Sense & Sensibility      3        0 austen
## 7 Sense & Sensibility      5        0 1811
## 8 Sense & Sensibility     10        1 chapter
## 9 Sense & Sensibility     10        1 1
## 10 Sense & Sensibility     13        1 the
## # ... with 725,045 more rows
```

Visualize the most common words in Jane Austen

```
tidy_austen %>%  
  anti_join(stop_words, by="word") %>%  
  count(word, sort=TRUE) %>%  
  top_n(15) %>%  
  mutate(word = reorder(word, n)) %>%  
  ggplot(aes(x=word, y=n)) +  
  geom_col() +  
  coord_flip()
```


Selecting by n



Visualize the most common words by novel

```
tidy_austen %>%  
  anti_join(stop_words, by="word") %>%  
  count(book, word, sort=TRUE) %>%  
  mutate(word = factor(word, levels = rev(unique(word)))) %>%  
  group_by(book) %>%  
  top_n(15) %>%  
  ungroup() %>%  
  ggplot(aes(word, n, fill = book)) +  
  geom_col(show.legend = FALSE) +  
  labs(x = NULL, y = "n") +  
  facet_wrap(~book, ncol = 3, scales = "free") +  
  coord_flip()
```

Selecting by n



Analyze word frequency

When analyzing a corpus, we may want to identify words that are the most important to each document. However, looking at frequency alone may not help, since words that appear very commonly across all documents will be ranked highly, yet are not very important to any particular document.

Instead, we could give a weighting to words that are used most frequently in some documents, and not as much in others. The *inverse document frequency* (*idf*) does this by downweighting words that are very common across all documents in a corpus.

$$idf(\text{term}) = \ln \left(\frac{n_{\text{documents}}}{n_{\text{documents containing term}}} \right)$$

The *idf* can be combined with the *tf*, or term frequency (by multiplying them together) to calculate the *tf-idf*, which measures the importance of a term to a particular document in a corpus.

Although it has no basis in statistics or information theory, the *tf-idf* is a commonly used heuristic measure.

Calculate term frequency

```
austen_words <- tidy_austen %>%  
  count(book, word, sort=TRUE)  
  
austen_totals <- austen_words %>%  
  group_by(book) %>%  
  summarise(total = sum(n))  
  
austen_words <- austen_words %>%  
  left_join(austen_totals) %>%  
  mutate(tf = n / total)
```

```
## Joining, by = "book"
```

```
austen_words
```

```
## # A tibble: 40,379 x 5
```

##	book	word	n	total	tf
##	<fct>	<chr>	<int>	<int>	<dbl>
##	1 Mansfield Park	the	6206	160460	0.0387
##	2 Mansfield Park	to	5475	160460	0.0341
##	3 Mansfield Park	and	5438	160460	0.0339
##	4 Emma	to	5239	160996	0.0325
##	5 Emma	the	5201	160996	0.0323
##	6 Emma	and	4896	160996	0.0304
##	7 Mansfield Park	of	4778	160460	0.0298
##	8 Pride & Prejudice	the	4331	122204	0.0354
##	9 Emma	of	4291	160996	0.0267
##	10 Pride & Prejudice	to	4162	122204	0.0341
##	# ... with 40,369 more rows				

Calculate inverse document frequency

```
austen_words <- austen_words %>%  
  mutate(n_doc = n_distinct(book)) %>%  
  group_by(word) %>%  
  mutate(idf = log(n_doc / n()),  
         tf_idf = tf * idf) %>%  
  select(-n_doc, -total) %>%  
  ungroup()
```

```
austen_words
```

```
## # A tibble: 40,379 x 6
```

##	book	word	n	tf	idf	tf_idf
##	<fct>	<chr>	<int>	<dbl>	<dbl>	<dbl>
##	1 Mansfield Park	the	6206	0.0387	0	0
##	2 Mansfield Park	to	5475	0.0341	0	0
##	3 Mansfield Park	and	5438	0.0339	0	0
##	4 Emma	to	5239	0.0325	0	0
##	5 Emma	the	5201	0.0323	0	0
##	6 Emma	and	4896	0.0304	0	0
##	7 Mansfield Park	of	4778	0.0298	0	0
##	8 Pride & Prejudice	the	4331	0.0354	0	0
##	9 Emma	of	4291	0.0267	0	0
##	10 Pride & Prejudice	to	4162	0.0341	0	0
##	# ... with 40,369 more rows					

```
austen_words %>% arrange(desc(tf_idf))
```

```
## # A tibble: 40,379 x 6
```

##	book	word	n	tf	idf	tf_idf
##	<fct>	<chr>	<int>	<dbl>	<dbl>	<dbl>
##	1 Sense & Sensibility	elinor	623	0.00519	1.79	0.00931
##	2 Sense & Sensibility	marianne	492	0.00410	1.79	0.00735
##	3 Mansfield Park	crawford	493	0.00307	1.79	0.00551
##	4 Pride & Prejudice	darcy	373	0.00305	1.79	0.00547
##	5 Persuasion	elliot	254	0.00304	1.79	0.00544
##	6 Emma	emma	786	0.00488	1.10	0.00536
##	7 Northanger Abbey	tilney	196	0.00252	1.79	0.00452
##	8 Emma	weston	389	0.00242	1.79	0.00433
##	9 Pride & Prejudice	bennet	294	0.00241	1.79	0.00431
##	10 Persuasion	wentworth	191	0.00228	1.79	0.00409
##	... with 40,369 more rows					

Calculating tf-idf with `bind_tf_idf()`

Alternatively, the *tf-idf* for each term can be calculated easily by using the `bind_tf_idf()` function provided by the `tidytext` package.

```
bind_tf_idf(data, term, document, n)
```

- ▶ The first argument is the data
- ▶ The following arguments give the name of the columns that identify the term, document, and counts

```
austen_words_tf <- tidy_austen %>%  
  count(book, word, sort=TRUE) %>%  
  bind_tf_idf(word, book, n)
```

```
austen_words_tf %>% arrange(desc(tf_idf))
```

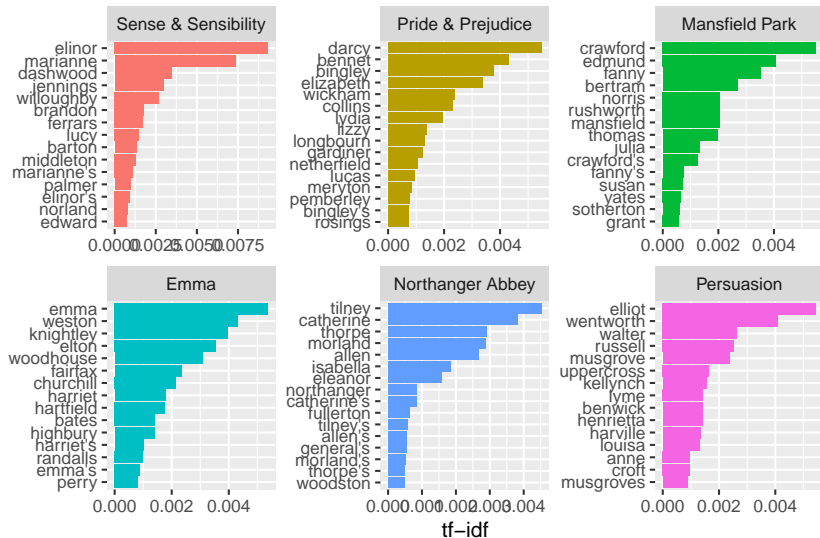
```
## # A tibble: 40,379 x 6
```

##	book	word	n	tf	idf	tf_idf
##	<fct>	<chr>	<int>	<dbl>	<dbl>	<dbl>
##	1 Sense & Sensibility	elinor	623	0.00519	1.79	0.00931
##	2 Sense & Sensibility	marianne	492	0.00410	1.79	0.00735
##	3 Mansfield Park	crawford	493	0.00307	1.79	0.00551
##	4 Pride & Prejudice	darcy	373	0.00305	1.79	0.00547
##	5 Persuasion	elliot	254	0.00304	1.79	0.00544
##	6 Emma	emma	786	0.00488	1.10	0.00536
##	7 Northanger Abbey	tilney	196	0.00252	1.79	0.00452
##	8 Emma	weston	389	0.00242	1.79	0.00433
##	9 Pride & Prejudice	bennet	294	0.00241	1.79	0.00431
##	10 Persuasion	wentworth	191	0.00228	1.79	0.00409
##	... with 40,369 more rows					

Visualizing the tf-idf of Jane Austen novels

```
austen_words %>%  
  arrange(desc(tf_idf)) %>%  
  mutate(word = factor(word, levels = rev(unique(word)))) %>%  
  group_by(book) %>%  
  top_n(15) %>%  
  ungroup() %>%  
  ggplot(aes(word, tf_idf, fill = book)) +  
  geom_col(show.legend = FALSE) +  
  labs(x = NULL, y = "tf-idf") +  
  facet_wrap(~book, ncol = 3, scales = "free") +  
  coord_flip()
```

Selecting by tf_idf



Analyzing n-grams

So far we have looked at the importance of individual words, but not the relationships between them. We can investigate relationships between words (e.g., which words tend to immediately follow others, or what words tend to co-occur in a corpus) by tokenizing on n-grams. We will focus on $n = 2$, or bigrams.

```
austen_bigrams <- austen_books() %>%  
  unnest_tokens(bigram, text, token = "ngrams", n = 2)
```

```
austen_bigrams
```

```
## # A tibble: 725,049 x 2
##   book                                bigram
##   <fct>                             <chr>
## 1 Sense & Sensibility sense and
## 2 Sense & Sensibility and sensibility
## 3 Sense & Sensibility sensibility by
## 4 Sense & Sensibility by jane
## 5 Sense & Sensibility jane austen
## 6 Sense & Sensibility austen 1811
## 7 Sense & Sensibility 1811 chapter
## 8 Sense & Sensibility chapter 1
## 9 Sense & Sensibility 1 the
## 10 Sense & Sensibility the family
## # ... with 725,039 more rows
```

Check most common bigrams in Jane Austen

```
austen_bigrams %>% count(bigram, sort=TRUE)
```

```
## # A tibble: 211,236 x 2
##   bigram      n
##   <chr>    <int>
## 1 of the    3017
## 2 to be    2787
## 3 in the   2368
## 4 it was   1781
## 5 i am     1545
## 6 she had  1472
## 7 of her   1445
## 8 to the   1387
## 9 she was  1377
## 10 had been 1299
## # ... with 211,226 more rows
```

Separate bigrams and remove stop words

```
austen_bigrams <- austen_bigrams %>%  
  separate(bigram, c("word1", "word2"), sep = " ")
```

```
austen_bigrams <- austen_bigrams %>%  
  filter(!word1 %in% stop_words$word) %>%  
  filter(!word2 %in% stop_words$word)
```


Check most common bigrams again

```
austen_bigrams %>% count(word1, word2, sort=TRUE)
```

```
## # A tibble: 33,421 x 3
##   word1    word2      n
##   <chr>   <chr>   <int>
## 1 sir      thomas    287
## 2 miss     crawford  215
## 3 captain wentworth  170
## 4 miss     woodhouse 162
## 5 frank    churchill 132
## 6 lady     russell   118
## 7 lady     bertram   114
## 8 sir      walter    113
## 9 miss     fairfax   109
## 10 colonel brandon   108
## # ... with 33,411 more rows
```

Check most common use of “miss”

```
austen_bigrams %>%  
  filter(word1 == "miss") %>%  
  count(word1, word2, sort=TRUE)
```

```
## # A tibble: 121 x 3  
##   word1 word2      n  
##   <chr> <chr>   <int>  
## 1 miss  crawford  215  
## 2 miss  woodhouse  162  
## 3 miss  fairfax   109  
## 4 miss  bates     103  
## 5 miss  tilney     82  
## 6 miss  bingley    72  
## 7 miss  dashwood   65  
## 8 miss  bennet     60  
## 9 miss  morland    55  
## 10 miss smith     52  
## # ... with 111 more rows
```

Calculate tf-idf for bigrams

```
austen_bigrams_tf <- austen_bigrams %>%  
  unite(bigram, word1, word2, sep = " ") %>%  
  count(book, bigram) %>%  
  bind_tf_idf(bigram, book, n) %>%  
  arrange(desc(tf_idf))
```

austen_bigrams_tf

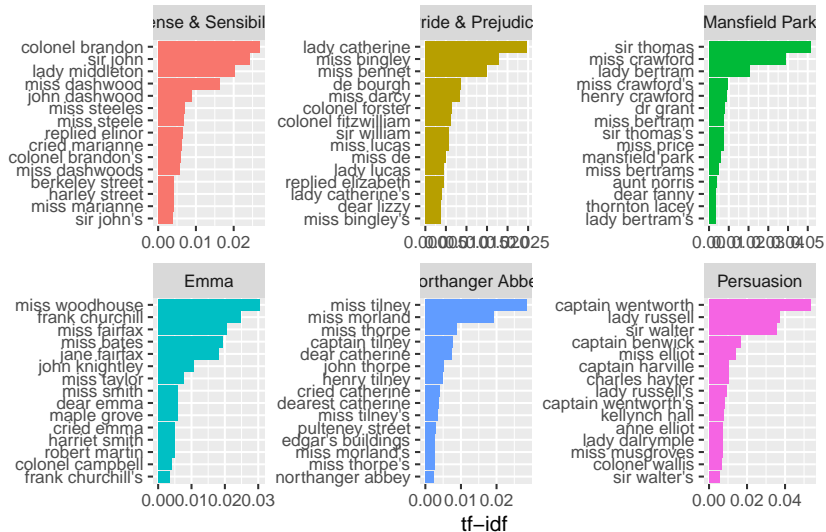
A tibble: 36,217 x 6

##	book	bigram	n	tf	idf	tf_idf
##	<fct>	<chr>	<int>	<dbl>	<dbl>	<dbl>
##	1 Persuasion	captain went~	170	0.0299	1.79	0.0535
##	2 Mansfield Park	sir thomas	287	0.0287	1.79	0.0515
##	3 Mansfield Park	miss crawford	215	0.0215	1.79	0.0386
##	4 Persuasion	lady russell	118	0.0207	1.79	0.0371
##	5 Persuasion	sir walter	113	0.0198	1.79	0.0356
##	6 Emma	miss woodhou~	162	0.0170	1.79	0.0305
##	7 Northanger Abbey	miss tilney	82	0.0159	1.79	0.0286
##	8 Sense & Sensibi~	colonel bran~	108	0.0150	1.79	0.0269
##	9 Emma	frank church~	132	0.0139	1.79	0.0248
##	10 Pride & Prejudi~	lady catheri~	100	0.0138	1.79	0.0247
##	# ... with 36,207 more rows					

Visualize bigrams by tf-idf

```
austen_bigrams_tf %>%  
  arrange(desc(tf_idf)) %>%  
  mutate(bigram = factor(bigram, levels = rev(unique(bigram))))  
  group_by(book) %>%  
  top_n(15) %>%  
  ungroup() %>%  
  ggplot(aes(bigram, tf_idf, fill = book)) +  
  geom_col(show.legend = FALSE) +  
  labs(x = NULL, y = "tf-idf") +  
  facet_wrap(~book, ncol = 3, scales = "free") +  
  coord_flip()
```

Selecting by tf_idf



Analyzing bigrams with graphs

Rather than visualize the most common bigrams by their counts or term frequencies, we can instead visualize them as a graph. Here, we use “graph” not in the sense of visualization, but to refer to a network of connected nodes.

A graph can be constructed from a tidy object by considering the *edges* between nodes as the observations (rows), and the variables (columns) are:

- ▶ **from** the node that the edge is coming from
- ▶ **to** the node that the edge is going toward
- ▶ **weight** the weight of the edge

In our case, each bigram represents an edge in the network, each word is a node, and the weights are the bigram counts.

We can use the `igraph` package to construct the graph.

Creating a graph from bigrams

```
library(igraph)
austen_graph <- austen_bigrams %>%
  count(word1, word2, sort=TRUE) %>%
  filter(n > 20) %>%
  graph_from_data_frame()

austen_graph
```

```
## IGRAPH 77e17f6 DN-- 91 77 --
## + attr: name (v/c), n (e/n)
## + edges from 77e17f6 (vertex names):
## [1] sir      ->thomas    miss      ->crawford
## [3] captain ->wentworth miss      ->woodhouse
## [5] frank   ->churchill lady      ->russell
## [7] lady    ->bertram   sir       ->walter
## [9] miss    ->fairfax   colonel   ->brandon
## [11] miss    ->bates     lady      ->catherine
## [13] sir     ->john      jane      ->fairfax
## [15] miss    ->tilney    lady      ->middleton
## + ... omitted several edges
```


Visualizing n-gram graphs

We can use the `ggraph` package to visualize the graph of bigrams that we just made.

```
library(ggraph)
```

```
set.seed(1)
ggraph(austen_graph) +
  geom_edge_link() +
  geom_node_point() +
  geom_node_text(aes(label = name), vjust = 1, hjust = 1)
```

```
## Using `nicely` as default layout
```

